Measurement checks — HPT (Czech data)
DIF & multi-group CFA by ideological attitude (ideological contamination tests)

HPT and Extremism project

2025-12-09

Contents

1 Purpose & plan

We assess whether *ideological attitude groups* respond differently to the HPT items **even at equal underlying HPT competence**. Concretely:

- Create **low/high ideology** groups from FR-LF-mini and KSA-3 (see codebook).

- Run **item-level DIF** (ordinal logistic) for each HPT item.
- Run **multi-group CFA** (configural → metric → scalar).
- **Interpretation:** If we find pervasive DIF and/or scalar non-invariance, this supports the PCI RR H1 that HPT scores can be ideologically contaminated.
- Background on HPT instrument and typical factor solutions:

2 Setup

```r
options(width = 120)
library(dplyr)
library(tidyr)
library(ggplot2)
library(psych)
library(knitr)
library(stringr)
library(janitor)
library(difR)          # DIF for ordinal items
library(lavaan)        # CFA / invariance
library(semTools)      # measurementInvariance, fitMeasures helpers
library(car)           # recode
library(mirt)
```

3 Data

```r
# Load the dataset created in 00_data-preparation
load("normalised_responses.RData")
stopifnot(exists("normalised_responses"))
dat_raw <- normalised_responses
```

We use variables as defined in the **codebook** (metadata; KN1-KN6; POP1-POP3; ROA1-ROA3; CONT1-CONT3; FR-LF mini RD1-RD3 & NS1-NS3; KSA-3 A1-A3, U1-U3, K1-K3; SDR1-SDR5).

3.1 Scoring & grouping

- **HPT items** are 1-4. To align a single "more HPT-competent" direction, we reverse only the *presentism* items (POP1–POP3): `POP*_rev = 5 - POP*`. CONT and ROA are used as-is (higher = better contextual/agent-aware reasoning per instrument notes).
- **Ideology composite**: FR-LF-mini (RD1-3 + NS1-3) and KSA-3 (9 items). We z-score the two scale means and average them → **IDEO_Z**. *Low =* bottom 33%, *High* = top 33% (middle third excluded to sharpen contrasts). FR-LF dimensions/logic per Leipzig scale.
- **Controls**: prior knowledge (sum KN1-KN6), social desirability (SDR1-SDR5; note SDR2-SDR4 already reversed in your file).

```r
# Select item blocks
hpt_items <- c(paste0("POP",1:3), paste0("ROA",1:3), paste0("CONT",1:3))
frlf_items <- c(paste0("RD",1:3), paste0("NS",1:3))
ksa_items  <- c(paste0("A",1:3), paste0("U",1:3), paste0("K",1:3))
kn_items   <- paste0("KN",1:6)
sdr_items  <- paste0("SDR",1:5)

dat <- dat_raw %>%
  mutate(across(all_of(hpt_items), as.numeric),
         across(all_of(frlf_items), as.numeric),
         across(all_of(ksa_items), as.numeric),
         across(all_of(kn_items), as.numeric),
         across(all_of(sdr_items), as.numeric)) %>%
  # Reverse presentism items so higher = better HPT
  mutate(across(all_of(paste0("POP",1:3)), ~ 5 - .x, .names = "{.col}_rev")) %>%
  # Scale scores
  rowwise() %>%
  mutate(
    HPT_total = mean(c_across(c(paste0("POP",1:3,"_rev"), ROA1:ROA3, CONT1:CONT3)), na.rm = TRUE),
    HPT_CONT  = mean(c_across(CONT1:CONT3), na.rm = TRUE),
    HPT_ROA   = mean(c_across(ROA1:ROA3), na.rm = TRUE),
    HPT_POPR  = mean(c_across(paste0("POP",1:3,"_rev")), na.rm = TRUE),
    FRLF_mean = mean(c_across(all_of(frlf_items)), na.rm = TRUE),
    KSA_mean  = mean(c_across(all_of(ksa_items)), na.rm = TRUE),
    KN_sum    = sum(c_across(all_of(kn_items)), na.rm = TRUE),
    SDR_mean  = mean(c_across(all_of(sdr_items)), na.rm = TRUE)
  ) %>%
  ungroup() %>%
  mutate(
    FRLF_z = as.numeric(scale(FRLF_mean)),
    KSA_z  = as.numeric(scale(KSA_mean)),
    IDEO_Z = (FRLF_z + KSA_z) / 2
  )
```

```
# define groups by tertiles
qs <- quantile(dat$IDEO_Z, probs = c(.3334, .6666), na.rm = TRUE)
dat <- dat %>%
  mutate(
    ideology_group = case_when(
      IDEO_Z <= qs[1] ~ "Low",
      IDEO_Z >= qs[2] ~ "High",
      TRUE ~ "Mid"
    )
  )


tab <- dat %>% count(ideology_group)
kable(tab, caption = "Group sizes (Low/High ideology; Mid excluded from group-wise tests)")
```

Table 1: Group sizes (Low/High ideology; Mid excluded from group-wise tests)

ideology_group	n
High	53
Low	53
Mid	50

Interpretation note. We focus on *Low* vs *High* ideology to maximize contrast and statistical power for DIF/MG-CFA. The *Mid* group is retained for descriptive context but excluded from two-group invariance tests.

4 Descriptives (checks)

```
desc_tbl <- dat %>%
  group_by(ideology_group) %>%
  summarise(n = n(),
            HPT_total = mean(HPT_total, na.rm = TRUE),
            KN_sum    = mean(KN_sum, na.rm = TRUE),
            SDR_mean  = mean(SDR_mean, na.rm = TRUE)) %>%
  arrange(match(ideology_group, c("Low","Mid","High")))
kable(desc_tbl, digits = 2, caption = "Descriptives by ideology group (means)")
```

Table 2: Descriptives by ideology group (means)

ideology_group	n	HPT_total	KN_sum	SDR_mean
Low	53	2.89	3.66	3.12
Mid	50	2.93	3.22	3.13
High	53	2.75	3.00	2.90

5 DIF analysis (ordinal logistic, item-by-item)

Goal. At *equal HPT ability*, do Low/High ideology groups respond differently to specific items? We use *matching on total HPT (item-rest)* and test both uniform and non-uniform DIF per item (= .01, Bonferroni adjusted). This targets the *"same trait, different item functioning"* mechanism of ideological contamination.

```r
## Keep only Low/High groups
anal <- dat[dat$ideology_group %in% c("Low","High"), , drop = FALSE]

## Ensure the HPT item list exists in the exact order we need
hpt_items <- c("POP1","POP2","POP3","ROA1","ROA2","ROA3","CONT1","CONT2","CONT3")

## Build item matrix with base subsetting (no tidy-eval)
stopifnot(all(hpt_items %in% names(anal)))
hpt_mat <- anal[, hpt_items, drop = FALSE]

## Coerce to numeric safely
for (j in seq_along(hpt_items)) {
  hpt_mat[[j]] <- suppressWarnings(as.numeric(hpt_mat[[j]]))
}

## Clean group factor
grp <- factor(anal$ideology_group, levels = c("Low","High"))

## Drop rows with < 2 non-missing item responses (mirt can't estimate for all-NA)
keep <- rowSums(!is.na(hpt_mat)) >= 2
hpt_mat <- hpt_mat[keep, , drop = FALSE]
grp     <- droplevels(grp[keep])

## Sanity checks
stopifnot(nrow(hpt_mat) == length(grp))
stopifnot(nlevels(grp) == 2)
print(table(grp))
```

```
## grp
##  Low High
##    53   53
```

```r
library(mirt)

# 1. Fit the Base Model (Configural: constrained slopes & intercepts across groups)
#    We start with a constrained model because scheme="drop" will try to FREE parameters to see if fit improves.
mod_base <- multipleGroup(
  data       = hpt_mat,
  model      = 1,
  group      = grp,
  itemtype   = "graded",
  invariance = c("slopes", "intercepts", "free_means", "free_var")
)
```

```r
# 2. Robust Parameter Detection
#    Instead of guessing "d" or "a1", we ask the model what parameters exist.
#    graded models usually have 'a1' (slope) and 'd1', 'd2', 'd3'... (thresholds).
params_all <- mirt::mod2values(mod_base)$name
unique_pars <- unique(params_all)

# Identify Slopes: usually start with 'a' (e.g., a1)
pars_slope <- grep("^a", unique_pars, value = TRUE)

# Identify Thresholds: usually start with 'd' (e.g., d1, d2, d3)
# Note: We ensure we don't pick up 'group' or other metadata parameters
pars_thr <- grep("^d\\d+$", unique_pars, value = TRUE)

if (length(pars_slope) == 0 || length(pars_thr) == 0) {
  stop("Could not auto-detect parameter names. Please check mod2values(mod_base).")
}

pars_to_test <- c(pars_slope, pars_thr)

# 3. Run DIF
dif_out <- DIF(
  mod_base,
  which.par  = pars_to_test,
  scheme     = "drop",
  items2test = colnames(hpt_mat),
  p.adjust   = "bonferroni",
```

```r
    verbose      = FALSE
)

# 4. Tidy Output (Dynamic Aggregation)
#    DIF() returns columns like p.a1, p.d1, p.d2. We need to aggregate them.
res_tbl <- as.data.frame(dif_out)
res_tbl$Item <- rownames(res_tbl)

# Helper to get min p-value across a set of columns (e.g., all d-parameters)
get_min_p <- function(df, prefix_list) {
  # Find columns in df matching p.prefix (e.g., p.d1, p.d2)
  cols <- unlist(lapply(prefix_list, function(p) grep(paste0("^p\\.", p, "$"), names(df), value=TRUE)))
  if (length(cols) == 0) return(rep(NA, nrow(df)))

  # Return min p-value across these columns for each row
  apply(df[, cols, drop=FALSE], 1, min, na.rm = TRUE)
}

alpha <- 0.01

res_tbl <- res_tbl %>%
  mutate(
    # Aggregate specific params into broad categories
    p_nonuniform = get_min_p(cur_data(), pars_slope), # Min p of a1...
    p_uniform    = get_min_p(cur_data(), pars_thr),   # Min p of d1, d2...

    # Flagging logic
    Flag_nonuniform = ifelse(!is.na(p_nonuniform) & p_nonuniform < alpha, "YES", "no"),
    Flag_uniform    = ifelse(!is.na(p_uniform)    & p_uniform    < alpha, "YES", "no")
  ) %>%
  select(Item, p_nonuniform, p_uniform, Flag_nonuniform, Flag_uniform)

knitr::kable(
  res_tbl, digits = 4,
  caption = "DIF per item (mirt; graded). Non-uniform = Slope (a1); Uniform = Any Threshold (d1-dK). Bonferroni  = 0.01."
)
```

Table 3: DIF per item (mirt; graded). Non-uniform = Slope (a1); Uniform = Any Threshold (d1-dK). Bonferroni = 0.01.

	Item	p_nonuniform	p_uniform	Flag_nonuniform	Flag_uniform
POP1	POP1	NA	NA	no	no

	Item	p_nonuniform	p_uniform	Flag_nonuniform	Flag_uniform
POP2	POP2	NA	NA	no	no
POP3	POP3	NA	NA	no	no
ROA1	ROA1	NA	NA	no	no
ROA2	ROA2	NA	NA	no	no
ROA3	ROA3	NA	NA	no	no
CONT1	CONT1	NA	NA	no	no
CONT2	CONT2	NA	NA	no	no
CONT3	CONT3	NA	NA	no	no

How to read it (unchanged conceptually):

- **Uniform DIF (thresholds/d)** → one group consistently picks higher/lower categories across the ability range.
- **Non-uniform DIF (slopes/a1)** → the group difference changes with latent ability (interaction).

```r
flagged <- with(res_tbl, Item[Flag_uniform == "YES" | Flag_nonuniform == "YES"])
if (length(flagged) > 0) {
  which_item <- which(colnames(hpt_mat) == flagged[1])
  plot(mod_base, type = "trace", which.items = which_item,
       facet_items = FALSE, groups = levels(grp))
} else {
  # CHANGED: " " to "alpha" to prevent LaTeX encoding error
  plot.new(); text(0.5, 0.5, "No DIF-flagged item at alpha = 0.01.")
}
```

No DIF−flagged item at alpha = 0.01.

6 Multi-group CFA (configural → metric → scalar)

Model. Prior studies often find POP and CONT forming one factor and ROA another. To remain conservative and close to your codebook structure, we specify a **three-factor model** (POP_rev, ROA, CONT) with POP items reversed so that all loadings point to *more HPT-congruent* responses. We then test invariance across Low vs High ideology groups.

```r
# Build analysis frame with reversed POP and intact ROA/CONT
cfad <- dat %>%
  filter(ideology_group %in% c("Low","High")) %>%
  transmute(
    ideology_group,
    POP1 = 5 - !!sym("POP1"),
    POP2 = 5 - !!sym("POP2"),
    POP3 = 5 - !!sym("POP3"),
    ROA1, ROA2, ROA3,
```

```
    CONT1, CONT2, CONT3
  )
```

```r
model_3f <- '
  POP =~ POP1 + POP2 + POP3
  ROA =~ ROA1 + ROA2 + ROA3
  CONT =~ CONT1 + CONT2 + CONT3
  # Allow the three reasoning modes to correlate
'
```

```r
# We use standard lavaan syntax to run the 3 steps explicitly.
# This avoids the "lavTestLRT" error caused by the deprecated semTools wrapper.

ord_items <- colnames(cfad)[colnames(cfad) != "ideology_group"]

# 1. Configural: Same factor structure, no constraints
fit_conf <- cfa(model_3f, data = cfad, group = "ideology_group",
                ordered = ord_items, estimator = "WLSMV")

# 2. Metric (Weak): Constrain loadings
fit_metr <- cfa(model_3f, data = cfad, group = "ideology_group",
                ordered = ord_items, estimator = "WLSMV",
                group.equal = c("loadings"))

# 3. Scalar (Strong): Constrain loadings and thresholds
fit_scal <- cfa(model_3f, data = cfad, group = "ideology_group",
                ordered = ord_items, estimator = "WLSMV",
                group.equal = c("loadings", "thresholds"))

# Extract Fit Measures
get_fit <- function(fit) {
  fm <- fitMeasures(fit, c("chisq.scaled", "df.scaled", "pvalue.scaled",
                           "cfi.scaled", "rmsea.scaled", "srmr"))
  return(fm)
}

# Combine into a table
fits <- bind_rows(
  Configural = get_fit(fit_conf),
  Metric     = get_fit(fit_metr),
  Scalar     = get_fit(fit_scal)
) %>%
```

```
  mutate(Model = c("Configural", "Metric", "Scalar")) %>%
  select(Model, everything())

# Display Table
fits %>%
  mutate(across(where(is.numeric), round, 3)) %>%
  kable(caption = "MG-CFA fit indices by invariance level (WLSMV).")
```

Table 4: MG-CFA fit indices by invariance level (WLSMV).

Model	chisq.scaled	df.scaled	pvalue.scaled	cfi.scaled	rmsea.scaled	srmr
Configural	53.755	48	0.263	0.973	0.049	0.093
Metric	64.620	54	0.153	0.949	0.063	0.106
Scalar	68.821	63	0.287	0.972	0.043	0.096

```
# Calculate Deltas manually based on the table above
# Practical decision rules (Chen, 2007): ΔCFI  -.010 and ΔRMSEA  .015 support invariance.

deltas <- tibble(
  step = c("Configural -> Metric", "Metric -> Scalar"),

  # Note: CFI usually goes DOWN as constraints are added (so we look at change)
  dCFI = c(fits$cfi.scaled[2] - fits$cfi.scaled[1],
           fits$cfi.scaled[3] - fits$cfi.scaled[2]),

  # RMSEA usually goes UP (or stays same)
  dRMSEA = c(fits$rmsea.scaled[2] - fits$rmsea.scaled[1],
             fits$rmsea.scaled[3] - fits$rmsea.scaled[2])
)

# FIX: Use mutate + across to round only numeric columns
deltas %>%
  mutate(across(where(is.numeric), round, 3)) %>%
  kable(caption = "Delta fit (CFI, RMSEA) across steps.")
```

Table 5: Delta fit (CFI, RMSEA) across steps.

step	dCFI	dRMSEA
Configural -> Metric	-0.023	0.014
Metric -> Scalar	0.023	-0.020

How to read this.

- If **metric holds** (small ΔCFI/ΔRMSEA), the *loadings* are equivalent across Low/High: the latent traits have the same meaning.
- If **scalar fails** (large deltas), *thresholds/intercepts* differ $\rightarrow$ **latent means are not comparable**, consistent with item-level bias. This would **support H1** (ideological contamination).

7 (Optional) Two-factor robustness check

Some studies merge POP & CONT into one factor and keep ROA separate; if desired, run the same invariance ladder for a 2-factor model and compare conclusions.

```
model_2f <- '
  CTX =~ POP1 + POP2 + POP3 + CONT1 + CONT2 + CONT3
  ROA =~ ROA1 + ROA2 + ROA3
'
measurementInvariance(model_2f, data = cfad, group = "ideology_group",
                      estimator = "WLSMV",
                      ordered = colnames(cfad)[colnames(cfad)!="ideology_group"])
```

8 Interpretation & reporting

8.1 DIF summary (per-item bias)

```
kable(res_tbl, caption = "Repeat: DIF results to reference in text.")
```

Table 6: Repeat: DIF results to reference in text.

	Item	p_nonuniform	p_uniform	Flag_nonuniform	Flag_uniform
POP1	POP1	NA	NA	no	no

	Item	p_nonuniform	p_uniform	Flag_nonuniform	Flag_uniform
POP2	POP2	NA	NA	no	no
POP3	POP3	NA	NA	no	no
ROA1	ROA1	NA	NA	no	no
ROA2	ROA2	NA	NA	no	no
ROA3	ROA3	NA	NA	no	no
CONT1	CONT1	NA	NA	no	no
CONT2	CONT2	NA	NA	no	no
CONT3	CONT3	NA	NA	no	no

- **If several POP or CONT items show DIF** (esp. uniform DIF favoring *High* ideology), this indicates that **at the same overall HPT ability** the High-ideology group tends to endorse "context-fitting" options more strongly → **measurement bias** consistent with H1.
- **If DIF is absent or rare**, the items function similarly across ideology groups → **weak evidence for ideological contamination** at the item level.

8.2 MG-CFA summary (scale-level bias)

```
kable(fits %>% mutate(across(where(is.numeric), round, 3)),
      caption = "Repeat: MG-CFA fit to reference in text.")
```

Table 7: Repeat: MG-CFA fit to reference in text.

Model	chisq.scaled	df.scaled	pvalue.scaled	cfi.scaled	rmsea.scaled	srmr
Configural	53.755	48	0.263	0.973	0.049	0.093
Metric	64.620	54	0.153	0.949	0.063	0.106
Scalar	68.821	63	0.287	0.972	0.043	0.096

```
kable(deltas %>% mutate(across(where(is.numeric), round, 3)),
      caption = "Repeat: Delta fit (CFI, RMSEA) thresholds.")
```

Table 8: Repeat: Delta fit (CFI, RMSEA) thresholds.

step	dCFI	dRMSEA
Configural -> Metric	-0.023	0.014
Metric -> Scalar	0.023	-0.020

- **Configural fit acceptable + metric holds** but **scalar fails** → the construct is similar but item thresholds differ → **comparisons of group means are biased**, supporting H1.
- **Scalar holds** → the full measurement is invariant → **no scale-level contamination**; HPT score comparisons across ideology are tenable.

9   What this means for H1 (PCI RR)

- **Supports H1 (contamination)** if (a) multiple items show DIF and/or (b) scalar invariance fails. This means HPT performance may be **inflated when students' ideology aligns with the vignette's decisions/justifications** (the mechanism articulated in the PCI RR plan).
- **Weakens H1** if DIF is negligible and scalar invariance holds, i.e., **items and scale behave equivalently** across ideology groups.

10   Reproducibility appendix

```
sessionInfo()
```

```
## R version 4.4.2 (2024-10-31)
## Platform: x86_64-pc-linux-gnu
## Running under: Ubuntu 24.04.3 LTS
##
## Matrix products: default
## BLAS:   /usr/lib/x86_64-linux-gnu/blas/libblas.so.3.12.0
## LAPACK: /usr/lib/x86_64-linux-gnu/lapack/liblapack.so.3.12.0
##
## locale:
##  [1] LC_CTYPE=C.UTF-8           LC_NUMERIC=C              LC_TIME=cs_CZ.UTF-8        LC_COLLATE=C.UTF-8
##  [5] LC_MONETARY=cs_CZ.UTF-8    LC_MESSAGES=C.UTF-8       LC_PAPER=cs_CZ.UTF-8       LC_NAME=C
##  [9] LC_ADDRESS=C               LC_TELEPHONE=C            LC_MEASUREMENT=cs_CZ.UTF-8 LC_IDENTIFICATION=C
##
## time zone: Europe/Prague
## tzcode source: system (glibc)
##
## attached base packages:
## [1] stats4    stats     graphics  grDevices utils     datasets  methods   base
##
## other attached packages:
##  [1] mirt_1.45.1    lattice_0.22-5 car_3.1-3      carData_3.0-5 semTools_0.5-7 lavaan_0.6-20 difR_6.1.0
##  [8] janitor_2.2.1  stringr_1.5.1  knitr_1.50     psych_2.4.12  ggplot2_4.0.1  tidyr_1.3.1    dplyr_1.1.4
##
## loaded via a namespace (and not attached):
##   [1] RColorBrewer_1.1-3   rstudioapi_0.17.1    audio_0.1-11        shape_1.4.6.1        magrittr_2.0.3
##   [6] TH.data_1.1-4        estimability_1.5.1   farver_2.1.2         nloptr_2.2.1         rmarkdown_2.29
##  [11] fs_1.6.5             vctrs_0.6.5          minqa_1.2.8          htmltools_0.5.8.1    forcats_1.0.0
##  [16] haven_2.5.4          cellranger_1.1.0     Formula_1.2-5        dcurver_0.9.3        parallelly_1.45.1
##  [21] testthat_3.3.1       sandwich_3.1-1       emmeans_1.10.6       rootSolve_1.8.2.4    zoo_1.8-14
```

```
##  [26] lubridate_1.9.4       admisc_0.39            lifecycle_1.0.4       iterators_1.0.14     pkgconfig_2.0.3
##  [31] Matrix_1.7-1          R6_2.6.1              fastmap_1.2.0        rbibutils_2.3        future_1.68.0
##  [36] snakecase_0.11.1      digest_0.6.37        Exact_3.3            vegan_2.7-2          progressr_0.18.0
##  [41] timechange_0.3.0      abind_1.4-8         httr_1.4.7           mgcv_1.9-1           compiler_4.4.2
##  [46] proxy_0.4-27          withr_3.0.2         S7_0.2.1             R.utils_2.13.0       MASS_7.3-61
##  [51] sessioninfo_1.2.3     GPArotation_2024.3-1 permute_0.9-8       gld_2.6.7            tools_4.4.2
##  [56] pbivnorm_0.6.0        future.apply_1.20.0 clipr_0.8.0          R.oo_1.27.1          glue_1.8.0
##  [61] quadprog_1.5-8        nlme_3.1-166        grid_4.4.2           cluster_2.1.8        generics_0.1.3
##  [66] gtable_0.3.6          tzdb_0.5.0          R.methodsS3_1.8.2    class_7.3-22         data.table_1.17.8
##  [71] lmom_3.2             hms_1.1.3            Deriv_4.1.6          foreach_1.5.2        pillar_1.10.0
##  [76] splines_4.4.2        survival_3.7-0       tidyselect_1.2.1     pbapply_1.7-4        reformulas_0.4.1
##  [81] gridExtra_2.3        deltaPlotR_1.6       xfun_0.54            expm_1.0-0           brio_1.1.5
##  [86] stringi_1.8.4        VGAM_1.1-14          yaml_2.3.10          boot_1.3-31          evaluate_1.0.5
##  [91] codetools_0.2-20     beepr_2.0            msm_1.8.2            tibble_3.2.1         cli_3.6.5
##  [96] xtable_1.8-4         DescTools_0.99.60    Rdpack_2.6.4         Rcpp_1.0.13-1        readxl_1.4.3
## [101] globals_0.18.0       polycor_0.8-1        coda_0.19-4.1        parallel_4.4.2       readr_2.1.5
## [106] lme4_1.1-38          listenv_0.10.0       glmnet_4.1-10        mvtnorm_1.3-2        SimDesign_2.21
## [111] scales_1.4.0         e1071_1.7-16         purrr_1.1.0          rlang_1.1.6          multcomp_1.4-28
## [116] mnormt_2.1.1         ltm_1.2-0
```